
Do we need more bikes?

Project in Statistical Machine Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 In this project we develop, and study different statistical machine learning models
2 for predicting whether the number of available bikes at a given hour should be
3 increased, a project by the District Department of Transportation in Washington
4 D.C. The training data set consists of 1600 instances of hourly bike rentals, and
5 a test set of 400 instances. The models for prediction we have used are: *Logistic*
6 *regression*, *Discriminant methods: LDA, QDA, k- Nearest Neighbour*, and *Tree*
7 *Based Methods*. We have found that THE MODEL gives best prediction, with
8 accuracy ??????

9 1 Plan

10 1.1 From Intro

- 11 (i) Explore and preprocess data
- 12 (ii) try some or all classification methods, which are these?
 - 13 • Logistic Regression
 - 14 • Discriminant analysis: LDA, QDA
 - 15 • K-nearest neighbor
 - 16 • Tree-based methods: classification trees, random forests, bagging
 - 17 • Boosting
- 18 (iii) Which of these are to be "put in production"?

19 1.2 From Data analysis task

- 20 • Can any trend be seen comparing different hours, weeks, months?
- 21 • Is there any difference between weekdays and holidays?
- 22 • Is there any trend depending on the weather?

23 1.3 From Implementation of methods

24 Each group member should implement one family each, who did what shall be clear!

25 DNNs are encouraged to be implemented, do this if there is time. (DNN is not a thing a group member can claim as their family.)

26 Implement a naive version, let's do: *Always low_bike_demand*

28 1.3.1 What to do with each method

- 29 1. Implement the method (each person individually)
- 30 2. Tune hyper-parameters, discuss how this is done (each person individually)
- 31 3. Evaluate with for example cross-validation. Don't use E_{k-fold} (what is that?) (need to do
- 32 together)
- 33 4. (optional) Think about input features, are all relevant? (together)

34 Before training, unify pre-processing FOR ALL METHODS and choose ONE OR MULTIPLE

35 metrics to evaluate the model. (is it necessary to have the same for all?, is it beneficial?) Examples:

- 36 • accuracy
- 37 • f1-score
- 38 • recall
- 39 • precision

40 Use same test-train split for ALL MODELS

41 2 Introduction

42 Statistical machine learning is a subject that aims to build and train algorithms, that analyse large

43 amount of data, and make predictions for the future, which are computed by using established

44 statistical models, and tools from functional analysis. This is a project in supervised, statistical

45 machine learning, where several models were created, and trained, in order to analyse which one of

46 them gives best prediction for the project "Do we need more bikes", where we want to understand,

47 and predict if there is a high, or low demand of city bikes in the public transportation of Washington,

48 a project by the District Department of Transportation in Washington D.C..

49 The data set used for training our models, consist of 15 variables, containing quantitative/qualitative

50 data. We developed several models, and evaluated them with cross-validation, in order to understand

51 which algorithm gives the best prediction.

52 3 Theoretical Background

53 3.1 Mathematical Overview of the Models

54 3.1.1 Logistic Regression

55 The backbone of logistic regression is linear regression, i.e. finding the least-squares solution to an
56 equation system

$$X\theta = y \quad (1)$$

57 given by the normal equations

$$X^T X \theta = X^T y \quad (2)$$

58 where X is the training data matrix, θ is the coefficient vector and b is the training output. The
59 parameter vector is then used in the sigmoid function:

$$\sigma(z) = \frac{e^z}{1 + e^z} : \mathbb{R} \rightarrow [0, 1], \quad (3)$$

$$z = x^T \theta, \quad (4)$$

60 where x is the testing input. This gives a statistical interpretation of the input vector. In the case of a
61 binary True/False classification, the value of the sigmoid function then determines the class.

62 3.1.2 Random forest

63 The random forest method is based upon decision trees, i.e. dividing the data point into binary
64 groups based on Gini-impurity, entropy or classification error, Gini being the most common. These
65 divisions are then used to create a binary tree shown in figure ??Tree) and where the leaf-nodes are
66 used to classify the target variables based on the input. As of itself the decision tree tends to have
67 unsatisfying results which leads to methods like random forest that boost its accuracy.

68 3.1.3 Non-parametric method: k-Nearest Neighbour

69 *k-Nearest Neighbour* (k -NN) is a distance based method that takes a k amount of points from the
70 training data set, called *neighbours*, computes the distance between them, then assumes that the
71 predicted value $\hat{y}(x_*)$ follows the trend of the k -nearest neighbours. Since k -NN uses the training
72 data explicitly it is also called a *nonparametric* method.

73 The k -NN method can be divided into several subcategories, inter alia *classification* k -NN method,
74 *regression* k -NN method. In this project, we are using the classification method, since we are trying
75 to predict in which of the two classes low, or high demand, the given, and predicted data points
76 belong.

77 The classification k -NN algorithm evaluates $\hat{y}(x_*)$ by computing the most frequently occurring class
78 among the k nearest neighbours. Here, we try to identify whether a data point belongs to the high
79 demand-class. Denote c = high demand class. For simplicity, assume Euclidean distance. Then

$$\hat{y}(x_*) = \arg \max_c \sum_{n \in \mathbb{N}} \chi_{(y_n=c)},$$

80 where y_i is the class of the nearest neighbour, χ is the characteristic function

$$\chi_{(y_i=c)} = \begin{cases} 1 & \text{if } y_n = c, \\ 0 & \text{otherwise.} \end{cases}$$

81 It is very common to use a weighted sum to predict the next value, i.e.

$$\hat{y}(x_*) = \arg \max_c \sum_{n \in \mathbb{N}} \frac{\chi_{(y_n=c)}}{d(x, x_n)},$$

82 where d is the standard Euclidean metric, computing the distance between an input x , and a neighbour
83 x_n .

84 When using this model it is important to choose an optimal k -value. There are several tests for this,
85 here we implement *uniform weighting*, and *distance weighting*. The first algorithm creates a k -NN

model for each new $k \in [1, 500]$, and trains the model with uniform weights, i.e. the contribution of all neighbours is equal. Similarly, the latter trains a k -NN classifier for each $k \in [1, 500]$, with the difference that it uses distance based weighting, i.e. closer neighbours have greater influence. After testing different upper boundaries for k , the two models gave good results in the interval $[1, 500]$, see Figure 1. From the figures, we can see that the second test gives a better value for k , since the plot follows smoother trend, in comparison to the uniform weighting test, which makes it easier to identify an optimal k value ($k = 120$). Moreover, the distance weighting algorithm is providing results for larger values of k , that is for $k \in [1, 400]$ before the curve converges, while the uniform weighting algorithm converges earlier, when $k = 120$. This means that for large k , both test algorithms make prediction based on the most common class in the data set, instead of making prediction based on the behaviour of the neighbours. Thus for sufficiently large k , for any given data point, the model will consider unnecessarily large amount of neighbours, and the prediction will be evaluated to belong to the most frequent class. Since the distance weighting has a larger range of k -value, it should be more trustworthy.

When $k = 120$, the accuracy of the model is 92%.




Figure 1: Test for choosing an optimal k -value.

3.1.4 Discriminant analysis: LDA and QDA

Linear Discriminant Analysis is a generative model, which means it is a model that's creating and using a probability distribution $P(\mathbf{x}, y)$ to create an estimation for the probability $P(y = m|\mathbf{x})$ using Bayes theorem.

Bayes theorem is:

$$p(y|\mathbf{x}) = \frac{p(y, \mathbf{x})}{p(\mathbf{x})} = \frac{p(y)p(\mathbf{x}|y)}{\int_y p(y, \mathbf{x})}$$

For the discrete version it is obtained:

$$p(y = m|\mathbf{x}) = \frac{p(y = m)p(\mathbf{x}|y = m)}{\sum_{m=1}^M p(y = m)p(\mathbf{x}|y = m)}$$

For this form of the equation to be useful, it is necessary to obtain an accurate estimation of $p(y = m)$ and $p(\mathbf{x}|y = m)$ for all classes m .

In LDA, $p(y = m)$ is estimated by counting the percentage of data points (in the training data) being in each of the classes and using that percentage as the probability of a data point being in that class. In mathematical terms:

$$p(y = m) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}\{y_i = m\} = \frac{n_m}{n}$$

To estimate the probability distribution $p(\mathbf{x}|y = m)$, a multi-dimensional gaussian distribution is used:

$$\mathcal{N}(\mathbf{x}|\mu, \Sigma) = \frac{1}{(2\pi)^{d/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$

Where \mathbf{x} is the d-dimensional data point, μ is the (d-dimensional) mean of the random variable. Σ is the symmetric, positive definite covariance matrix defined by:

$$\Sigma = \frac{1}{n - M} \sum_{m=1}^M \sum_{i:y_i=m} (\mathbf{x}_i - \mu_m)(\mathbf{x}_i - \mu_m)^T$$

Using these estimations results in an expression for the quantity $p(y = m|\mathbf{x})\forall m$. LDA then uses maximum likelihood to categorize an input \mathbf{x} into a class m .

Quadratic discriminant analysis (QDA) is heavily based on LDA with the sole difference being how the covariance matrix Σ is created. In LDA, the covariance matrix is assumed to be the same for data in each and every class. In QDA however, the covariance matrix is calculated for each class as follows:

$$\Sigma_m = \frac{1}{n_m - 1} \sum_{i:y_i=m} (\mathbf{x}_i - \mu_m)(\mathbf{x}_i - \mu_m)^T$$

One thing to note about LDA and QDA is that the use of a multi-variable gaussian distribution benefits normally distributed variables. In this project however, there is a dependence on positive definite values which are not normally distributed by nature. This is an issue when using QDA since in the class of *high_bike_demand*, all data points have a snow depth of 0 and has hence no variance. This results in this class having an undefined inverse for the covariance matrix. The solution used was to exclude this variable from this model.

3.2 Input Data Modification

By plotting the data and analyzing the .csv file, some observations were made. The different inputs were then changed accordingly:


- *Kept as-is:* weekday, windspeed, visibility, temp
- *Modified:*
 - month - split into two inputs, one cosine and one sine part. This makes the new inputs linear and can follow the fluctuations of the year. The original input was discarded.
 - hour_of_day - split into three boolean variables: demand_day, demand_evening, and demand_night, reflecting if the time was between 08-14, 15-19 or 20-07 respectively. This was done because plotting the data showed three different plateaus of demand for the different time intervals. The original input was discarded.
 - snowdepth, precip were transformed into booleans, reflecting if it was raining or if there was snow on the ground or not. This was done as there were no times where demand was high when it was raining or when there was snow on the ground.
- *Removed:* cloudcover, day_of_week, snow, dew, holiday, summertime. These were removed due to being redundant (e.g. summertime), not showing a clear trend (e.g. cloudcover), giving a worse score when used, or all three (e.g. day_of_week).

4 Data Analysis

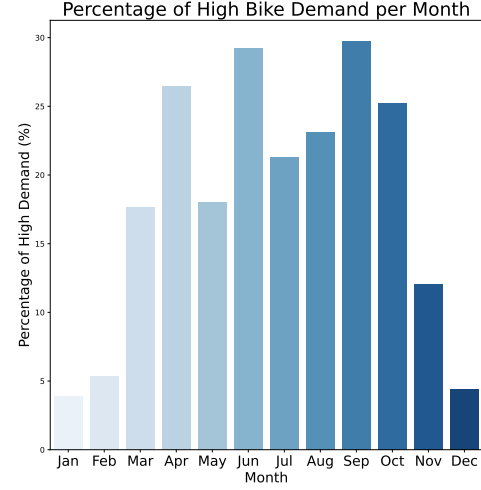
In the given data, there are some numerical and categorical features:

- *Numerical:* temp, dew, humidity, precip, snow, snowdepth, windspeed, cloudcover and visibility.
- *Categorical:* hour_of_day, day_of_week, month, holiday, weekday, summertime, and increase_stock


There are some trends seen in the data when it comes to time and weather. From figure 2, one can see a periodic relationship for the months, where there is a higher demand during the warmer months, loosely following a trigonometric curve. Over the week, the demand is rather stable, with a peak on the weekend, especially Saturdays.



(a) Demand per day of week.



(b) Demand per month.



(c) Demand per hour of day.

Figure 2: Bike demand vs. day of week and month.

Looking at the weather (figure 3); if there is rain or if there is snow on the ground, there is close to always low demand. Cloudcover did not make a big impact, which is also intuitive, as a cloudy day does not make biking more difficult. Dew point also does not have a clear trend, while humidity however has a clear trend downwards as the humidity increases. Temperature had a more clear impact, where more people wanted to bike the warmer it got.

The overall trend is that about one eighth of observations correspond to a high bike demand. During the night, or in bad weather, the demand is (intuitively) low. But during rush hour (figure 2c), the demand is very high, and should probably be increased in order to minimize excessive CO₂ emissions.

5 Result

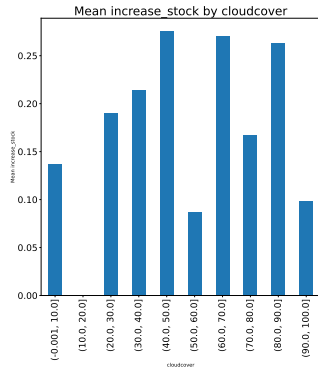
The method used to evaluate the different models where simply chosen to be the accuracy defined by:

$$\text{accuracy} = \frac{n_{\text{correct}}}{n_{\text{tot}}}$$

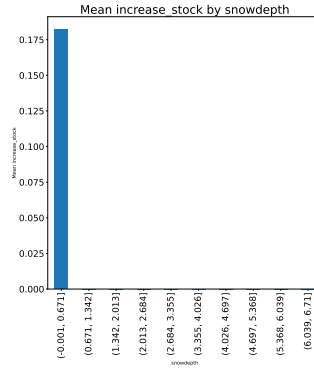
The different models were tested and the accuracy where:

Accuracy of the models

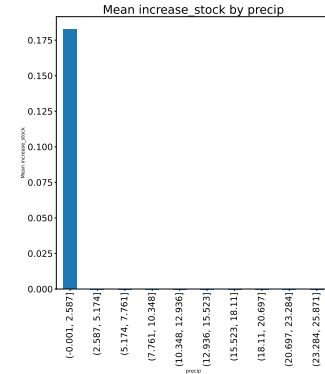
Model	Accuracy
LDA	86%
QDA	86%




(a) Demand per cloudcover (percentage).




(b) Demand per day of week.




(c) Demand per day of week.



(d) Demand per dew point ($^{\circ}\text{C}$).



(e) Demand per temperature ($^{\circ}$).



(f) Demand per humidity level (percentage).

Figure 3: Bike demand vs. various weather parameters.

168 A Appendix

```

169 1 import pandas as pd
170 2 import numpy as np
171 3 from sklearn.model_selection import train_test_split
172 4 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
173 5 from sklearn.linear_model import LogisticRegression
174 6 from sklearn.metrics import accuracy_score
175 7 from sklearn.metrics import classification_report
176 8
177 9 df = pd.read_csv('training_data_vt2025.csv')
178 10
179 11 # modify the month to represent the periodicity that is observed in
180    data.
181 12 df['month_cos'] = np.cos(df['month']*2*np.pi/12)
182 13 df['month_sin'] = np.sin(df['month']*2*np.pi/12)
183 14
184 15 # time of day, replaced with 3 bool values: is_night, is_day and
185    is_evening,
186 16 # adding the new categories back in the end.
187 17 def categorize_demand(hour):
188 18     if 20 <= hour or 7 >= hour:
189 19         return 'night'
190 20     elif 8 <= hour <= 14:
191 21         return 'day'
192 22     elif 15 <= hour <= 19:
193 23         return 'evening'
194 24
195 25 df['time_of_day'] = df['hour_of_day'].apply(categorize_demand)
196 26 df_dummies = pd.get_dummies(df['time_of_day'], prefix='is', drop_first
197    =False)
198 27 df = pd.concat([df, df_dummies], axis=1)
199 28
200 29 # Create bool of snowdepth and percipitation
201 30 df['snowdepth_bool'] = df['snowdepth'].replace(0, False).astype(bool)
202 31 df['precip_bool'] = df['precip'].replace(0, False).astype(bool)
203 32
204 33 # Seperate training data from target
205 34 X=df[['#holiday',
206 35     'weekday',
207 36     '#summertime',
208 37     'temp',
209 38     '#dew',
210 39     '#humidity',
211 40     '#visibility',
212 41     '#windspeed',
213 42     '#month',
214 43     'month_cos',
215 44     'month_sin',
216 45     '#hour_of_day',
217 46     'is_day',
218 47     'is_evening',
219 48     'is_night',
220 49     '#hour_cos',
221 50     '#hour_sin',
222 51     'snowdepth_bool',
223 52     'precip_bool'
224 53     ]]
225 54
226 55 y=df['increase_stock']
227 56
228 57 # Split dataset into training and test sets
229 58 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
230    =0.2, random_state=42)
231 59

```



```

23250 # Apply Linear Discriminant Analysis (LDA)
23351 lda = LinearDiscriminantAnalysis(n_components=1)
23452 X_train_lda = lda.fit_transform(X_train, y_train)
23553 X_test_lda = lda.transform(X_test)
23654
23755 # Train a classifier (Logistic Regression)
23856 clf = LogisticRegression()
23957 clf.fit(X_train_lda, y_train)
24058
24159 # Make predictions
24260 y_pred = clf.predict(X_test_lda)
24361
24462 # Evaluate accuracy
24563 accuracy = accuracy_score(y_test, y_pred)
24664 print(f"Model Accuracy: {accuracy:.2f}")
24765
24866 print(classification_report(y_test, y_pred))

```

Listing 1: Code for LDA

```

249 1 import pandas as pd
250 2 import numpy as np
251 3 from sklearn.model_selection import train_test_split
252 4 from sklearn.discriminant_analysis import
253     QuadraticDiscriminantAnalysis
254 5 from sklearn.metrics import accuracy_score
255 6 from sklearn.metrics import classification_report
256 7
257 8 df = pd.read_csv('training_data_vt2025.csv')
258 9
25910 # modify the month to represent the periodicity that is observed in
260     data.
26111 df['month_cos'] = np.cos(df['month']*2*np.pi/12)
26212 df['month_sin'] = np.sin(df['month']*2*np.pi/12)
26313
26414 # time of day, replaced with 3 bool values: is_night, is_day and
265     is_evening,
26615 # adding the new categories back in the end.
26716 def categorize_demand(hour):
26817     if 20 <= hour or 7 >= hour:
26918         return 'night'
27019     elif 8 <= hour <= 14:
27120         return 'day'
27221     elif 15 <= hour <= 19:
27322         return 'evening'
27423
27524 df['time_of_day'] = df['hour_of_day'].apply(categorize_demand)
27625 df_dummies = pd.get_dummies(df['time_of_day'], prefix='is', drop_first
277     =False)
27826 df = pd.concat([df, df_dummies], axis=1)
27927
28028 # Create bool of snowdepth and percipitation
28129 df['snowdepth_bool'] = df['snowdepth'].where(df['snowdepth'] == 0, 1)
28230 df['precip_bool'] = df['precip'].where(df['precip'] == 0, 1)
28331
28432 # Seperate training data from target
28533 X=df[['#holiday',
28634     'weekday',
28735     '#summertime',
28836     'temp',
28937     '#dew',
29038     '#humidity',
29139     '#visibility',
29240     '#windspeed',
29341     '#month',

```

```

29412         'month_cos',
29513         'month_sin',
29614         #'hour_of_day',
29715         'is_day',
29816         'is_evening',
29917         'is_night',
30018         #'snowdepth_bool',
30119         'precip_bool'
30220     ]
30321
30422 y=df['increase_stock']
30523
30624 # Split dataset into training and test sets
30725 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
308         =0.2, random_state=42)
30926
31027 # Apply Quadratic Discriminant Analysis (QDA)
31128 qda = QuadraticDiscriminantAnalysis()
31229 X_train_lda = qda.fit(X_train, y_train)
31330
31431 # Make predictions
31532 y_pred = qda.predict(X_test)
31633
31734 # Evaluate accuracy
31835 accuracy = accuracy_score(y_test, y_pred)
31936 print(f"Model Accuracy: {accuracy:.2f}")
32037
32138 print(classification_report(y_test, y_pred))

```

Listing 2: Code for QDA